

Distributed AI Model Training (Simplified Simulation)

➤ Objective:

Implement a mini simulation of Federated Learning where multiple nodes train partial models on subsets of data and aggregate results.

➤ Requirements:

- Dataset: Use the Iris dataset (from scikit-learn). Split into 3–4 subsets to simulate different nodes.
- Model Training: Each node trains a logistic regression model (use only scikit-learn).
- Aggregation: Implement a simple Federated Averaging (FedAvg) function to combine model weights.
- Simulation Features: Randomly drop one node in some rounds (skip its update). When the node recovers, include it again in aggregation.
- Comparison: Train a centralized logistic regression (on full dataset) and compare accuracy vs distributed training.
- Logging: Print logs for each round: node updates, aggregation, node failures/recoveries.

➤ Optional (Bonus):

- Plot accuracy over training rounds.
- Simulate different data per node (non-iid split).

➤ Expected Deliverables:

- Python code simulating distributed nodes.
- FedAvg aggregation implementation.
- Console logs showing training + aggregation + node failures.
- Accuracy comparison (centralized vs distributed).

➤ Evaluation Criteria:

- Correctness of distributed simulation.
- Code clarity and readability.
- Handling of node failures.
- Accuracy comparison.